

Yifan Yuan

Email: yifan.yuan@intel.com

Website: <https://yifanyuan3.github.io>

Research Interests

- Networking hardware and system software for datacenter
- Hardware-software co-design for distributed/disaggregated systems acceleration

Education

- **University of Illinois at Urbana-Champaign** *August 2017 – May 2022*
 - M.S. (2019), Ph.D. (2022) in Computer Engineering
 - Advisor: Prof. Nam Sung Kim
 - Doctoral committee: Prof. Nam Sung Kim, Dr. Ren Wang, Prof. Radhika Mittal, Prof. Deming Chen
- **Zhejiang University** *September 2014 – June 2018*
 - B.E. in Electronic Information Engineering

Work Experience

- **Intel Labs** *July 2022 – Present*
 - Research Scientist at Networking Platforms Lab, Hillsboro, OR
- **Microsoft Research** *June 2020 – August 2020*
 - Research Intern at Systems Research Group, Redmond, WA
- **Intel Labs** *May 2019 – August 2019*
May 2018 – August 2018
 - Research Intern at Networking Platforms Lab, Hillsboro, OR

Publications

- **A Quantitative Analysis and Guidelines of Data Streaming Accelerator in Intel 4th Gen Xeon Scalable Processors**
R. Kuper, I. Jeong, **Y. Yuan**, R. Wang, N. Ranganathan, N. Rao, J. Hu, S. Kumar, P. Lantz, N. S. Kim
The ACM Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), 2024
Paper based on real Intel product and software ecosystem
- **Demystifying CXL Memory with Genuine CXL-Ready Systems and Devices**
Y. Sun, **Y. Yuan**, Z. Yu, R. Kuper, C. Song, J. Huang, H. Ji, S. Agarwal, J. Lou, I. Jeong, R. Wang, J. H. Ahn, T. Xu, N. S. Kim
The ACM/IEEE International Symposium on Microarchitecture (MICRO), 2023
- **STYX: Exploiting SmartNIC Capability to Reduce Datacenter Memory Tax**
H. Ji, Y. Sun, M. Mansi, **Y. Yuan**, J. Huang, R. Kuper, M. Swift, N. S. Kim
The USENIX Annual Technical Conference (ATC), 2023
- **RAMBDA: RDMA-driven Acceleration Framework for Memory-intensive us-scale Datacenter Applications**
Y. Yuan, J. Huang, Y. Sun, T. Wang, J. Nelson, D. Ports, Y. Wang, R. Wang, C. Tai, N. S. Kim
The IEEE International Symposium on High-Performance Computer Architecture (HPCA), 2023
- **IDIO: Network-Driven, Inbound Network Data Orchestration on Server Processors**
M. Alian, S. Agarwal, J. Shin, N. Patel, **Y. Yuan**, D. Kim, R. Wang, N. S. Kim
The ACM/IEEE International Symposium on Microarchitecture (MICRO), 2022
- **Unlocking the Power of Inline Floating-Point Operations on Programmable Switches**
Y. Yuan, O. Alama, J. Fei, J. Nelson, D. R. K. Ports, A. Sapio, M. Canini, N. S. Kim
The USENIX Symposium on Networked Systems Design and Implementation (NSDI), 2022
- **Don't Forget the I/O When Allocating Your LLC**
Y. Yuan, M. Alian, Y. Wang, R. Wang, I. Kurakin, C. Tai, N. S. Kim
The ACM/IEEE International Symposium on Computer Architecture (ISCA), 2021
Code merged into Intel official RDT (pqos) library
- **QEI: Query Acceleration Can be Generic and Efficient in the Cloud**
Y. Yuan, Y. Wang, R. Wang, R. Chowdhury, C. Tai, N. S. Kim
The IEEE International Symposium on High-Performance Computer Architecture (HPCA), 2021

- **Data Direct I/O Characterization for Future I/O System Exploration**

M. Alian, **Y. Yuan**, J. Zhang, R. Wang, M. Jung, N. S. Kim

The IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS), 2020

- **HALO: Accelerating Flow Classification for Scalable Packet Processing in NFV**

Y. Yuan, Y. Wang, R. Wang, J. Huang

The ACM/IEEE International Symposium on Computer Architecture (ISCA), 2019

- **Accelerating Distributed Reinforcement Learning with In-Switch Computing**

Y. Li, I. Liu, **Y. Yuan**, D. Chen, A. Schwing, J. Huang

The ACM/IEEE International Symposium on Computer Architecture (ISCA), 2019

- **Project Almanac: A Time-Traveling Solid-State Drive**

X. Wang, **Y. Yuan**, Y. Zhou, C. C. Coats, J. Huang

The ACM European Conference on Computer Systems (EuroSys), 2019

- **A Network-Centric Hardware/Algorithm Co-Design to Accelerate Distributed Training of Deep Neural Networks**

Y. Li, J. Park, M. Alian, **Y. Yuan**, Q. Zheng, P. Pan, R. Wang, A. Schwing, H. Esmaeilzadeh, N. S. Kim

The ACM/IEEE International Symposium on Microarchitecture (MICRO), 2018

Patents

- **Data Consistency and Durability over Distributed Persistent Memory Systems**

R. Wang, **Y. Yuan**, Y. Wang, T.-Y. C. Tai, T. Hurson

US Patent 11,709,774, granted Jul. 2023

- **Workload Scheduler for Memory Allocation**

Y. Wang, R. Wang, T.-Y. C. Tai, **Y. Yuan**, P. Pathak, S. Vedantham, C. Macnamara

US Patent App. 16/799,745, filed Feb. 2020

- **Offload of Data Lookup Operations**

R. Wang, A. J. Herdrich, T.-Y. C. Tai, Y. Wang, R. Kondapalli, A. Bachmutsky, **Y. Yuan**

US Patent 11,698,929, granted Jul. 2023

Professional Service

- **Tutorial Organizer and Presenter:** On-chip Accelerators in 4th Gen Intel® Xeon® Scalable Processors: Features, Performance, Use Cases, and Future! (ISCA'2023)

- **Program Committee:** ISCA'2024 (ERC), HPCA'2024, HPCA'2023 (ERC), EuroSys'2022 (shadow PC)

- **Reviewer:** IEEE Transactions on Parallel and Distributed Systems (TPDS, 2022), IEEE Computer Architecture Letter (CAL, 2022-2023)

Research Experience

- **Embracing Emerging CXL devices in Modern Datacenter**

2022 – Present

Intel Labs

CXL has been attracting much attention as the next generation of device interconnect standard, providing unique features such as memory expansion and cache coherence. We have been exploring both CXL memory devices and CXL accelerators based on real commodity hardware from different aspects, including both hardware functions enhancements and system software optimizations. The results have been published in **MICRO'23**.

- **Accelerator Design for Network/Application Dataplane Operations**

2018 – Present

UIUC and Intel Labs

Tackling the “datacenter tax” problem and the “killer microsecond” problem, we design accelerator architecture, programming models, and integration schemes to accelerate a wide range of fine-grained but costly operations in datacenter’s software stacks and applications. The results have been published in **HPCA'21** and **ISCA'19**.

- **I/O Subsystem Design and Optimization for Modern Server CPU**

2018 – 2021

UIUC and Intel Labs

High-speed I/O devices can exert significant pressure on the CPU’s cache/memory system. We study the I/O-host interaction behavior in the real system, and build realistic and accurate I/O subsystem models for gem5 simulator. We also propose multiple solutions in both real systems and simulation models to optimize the data transfer, notification, and interference in the I/O subsystem. The results have been published in **MICRO'22**, **ISCA'21** and **ISPASS'20**.

- **In-network Computing for Distributed ML Training Acceleration**

2017 – 2021

UIUC and Microsoft Research

Distributed ML training is notoriously time- and resource-consuming. We propose to leverage the networking devices, including NICs (for in-network gradient compression) and switches (for in-network gradient aggregation), to facilitate the inter-machine communication, which is the most expensive portion in distributed training. We also

explore the new potential for P4 programmable switch to process more complicated (floating-point) operations. The results have been published in ***NSDI'22***, ***ISCA'19***, and ***MICRO'18***.

Teaching Experience

- **ECE 411:** Computer Organization and Design (UIUC, SP 2021)

Skills and Techniques

- **Programming languages:** C/C++, Verilog HDL, VHDL, Python, P4, Shell script, LaTeX, Matlab, etc.
- **Development skills:** Unix/Linux, FPGA, DPDK, RDMA, programmable switch, CUDA, gem5 simulator, sniper simulator, etc.