

Yifan Yuan

Email: yifan3@gmail.com

Website: <https://yifanyuan3.github.io>

Research Interests

- Networking hardware and system software for datacenter
- Hardware-software co-design for distributed/disaggregated systems acceleration
- Modern heterogeneous memory system and software

Education

- **University of Illinois at Urbana-Champaign** *August 2017 – May 2022*
 - M.S. (2019), Ph.D. (2022) in Computer Engineering
 - Advisor: Prof. Nam Sung Kim
- **Zhejiang University** *September 2014 – June 2018*
 - B.E. in Electronic Information Engineering

Work Experience

- **Meta** *June 2024 – Present*
 - Senior Research Engineer at AI-Systems Co-Design Team, Menlo Park, CA
- **Intel Labs** *July 2022 – May 2024*
 - Research Scientist at Systems Architecture Lab, Santa Clara, CA
- **Microsoft Research** *June 2020 – August 2020*
 - Research Intern at Systems Research Group, Redmond, WA
- **Intel Labs** *May 2019 – August 2019*
May 2018 – August 2018
 - Research Intern at Networking Platforms Lab, Hillsboro, OR

Publications

- **DCPerf: An Open-Source, Battle-Tested Performance Benchmark Suite for Datacenter Workloads**
W. Su, A. Dhanotia, C. Torres, J. Gandhi, N. Gholkar, S. Kanaujia, M. Naumov, K. Subramanian, V. Andrei, **Y. Yuan**, C. Tang
The ACM/IEEE International Symposium on Computer Architecture (ISCA, industry track), 2025
Paper based on real experience at Meta
- **Dynamic Load Balancer in Intel Xeon Scalable Processor: Performance Analyses, Enhancements, and Guidelines**
J. Lou, S. Vanavasam, **Y. Yuan**, R. Wang, N. S. Kim
The ACM/IEEE International Symposium on Computer Architecture (ISCA), 2025
Paper based on real Intel product and software ecosystem
- **A4: Microarchitecture-Aware LLC Management for Datacenter Servers with Emerging I/O Devices**
H. Park, J. Lou, S. Lee, **Y. Yuan**, K. Park, Y. Son, I. Jeong, N. S. Kim
The ACM/IEEE International Symposium on Computer Architecture (ISCA), 2025
- **M5: Mastering page migration and memory management for CXL-based tiered memory systems**
Y. Sun, J. Kim, Z. Yu, J. Zhang, S. Chai, M. J. Kim, H. Nam, J. Park, E. Na, **Y. Yuan**, R. Wang, J. H. Ahn, T. Xu, N. S. Kim
The ACM Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), 2025
- **Scaling Persistent In-memory Key-Value Stores over Modern Tiered, Heterogeneous Memory Hierarchies**
M. Cai, J. Shen, **Y. Yuan**, Z. Qu, B. Ye
IEEE Transactions on Computers (TC), 2024
- **Demystifying a CXL Type-2 Device: A Heterogeneous Cooperative Computing Perspective**
H. Ji, S. Vanavasam, Y. Zhou, Q. Xia, J. Huang, **Y. Yuan**, R. Wang, P. Gupta, B. Chitlur, I. Jeong, N. S. Kim
The ACM/IEEE International Symposium on Microarchitecture (MICRO), 2024
The first CXL Type-2 device paper based on real CXL systems
- **Nomad: Non-Exclusive Memory Tiering via Transactional Page Migration**
L. Xiang, Z. Lin, W. Deng, H. Lu, J. Rao, **Y. Yuan**, R. Wang
The USENIX Symposium on Operating Systems Design and Implementation (OSDI), 2024

- **Intel Accelerator Ecosystem: An SoC-Oriented Perspective**
Y. Yuan, R. Wang, N. Ranganathan, N. Rao, S. Kumar, P. Lantz, V. Sanjeevan, J. Cabrera, A. Kwatra, R. Sankaran, I. Jeong, N. S. Kim
The ACM/IEEE International Symposium on Computer Architecture (ISCA, industry track), 2024
Paper based on real Intel product and software ecosystem
- **A Quantitative Analysis and Guidelines of Data Streaming Accelerator in Modern Intel Xeon Scalable Processors**
R. Kuper, I. Jeong, Y. Yuan, R. Wang, N. Ranganathan, N. Rao, J. Hu, S. Kumar, P. Lantz, N. S. Kim
The ACM Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS), 2024
Paper based on real Intel product and software ecosystem
- **BonsaiKV: Towards Fast, Scalable, and Persistent Key-Value Stores with Tiered, Heterogeneous Memory System**
M. Cai, J. Shen, Y. Yuan, Z. Qu, B. Ye
The International Conference on Very Large Databases (VLDB), 2024
- **Demystifying CXL Memory with Genuine CXL-Ready Systems and Devices**
Y. Sun, Y. Yuan, Z. Yu, R. Kuper, C. Song, J. Huang, H. Ji, S. Agarwal, J. Lou, I. Jeong, R. Wang, J. H. Ahn, T. Xu, N. S. Kim
The ACM/IEEE International Symposium on Microarchitecture (MICRO), 2023
The first CXL memory paper based on real CXL systems
Covered by Semiconductor Engineering, Intel Community Blog, Hacker News, The New Stack
- **STYX: Exploiting SmartNIC Capability to Reduce Datacenter Memory Tax**
H. Ji, Y. Sun, M. Mansi, Y. Yuan, J. Huang, R. Kuper, M. Swift, N. S. Kim
The USENIX Annual Technical Conference (ATC), 2023
- **RAMBDA: RDMA-driven Acceleration Framework for Memory-intensive us-scale Datacenter Applications**
Y. Yuan, J. Huang, Y. Sun, T. Wang, J. Nelson, D. Ports, Y. Wang, R. Wang, C. Tai, N. S. Kim
The IEEE International Symposium on High-Performance Computer Architecture (HPCA), 2023
- **LADIO: Leakage-Aware Direct I/O for I/O-Intensive Workloads**
I. Jeong, J. Lou, Y. Son, Y. Park, Y. Yuan, N. S. Kim
Computer Architecture Letters (CAL), 2023
- **IDIO: Network-Driven, Inbound Network Data Orchestration on Server Processors**
M. Alian, S. Agarwal, J. Shin, N. Patel, Y. Yuan, D. Kim, R. Wang, N. S. Kim
The ACM/IEEE International Symposium on Microarchitecture (MICRO), 2022
- **Unlocking the Power of Inline Floating-Point Operations on Programmable Switches**
Y. Yuan, O. Alama, J. Fei, J. Nelson, D. R. K. Ports, A. Sapio, M. Canini, N. S. Kim
The USENIX Symposium on Networked Systems Design and Implementation (NSDI), 2022
- **Don't Forget the I/O When Allocating Your LLC**
Y. Yuan, M. Alian, Y. Wang, R. Wang, I. Kurakin, C. Tai, N. S. Kim
The ACM/IEEE International Symposium on Computer Architecture (ISCA), 2021
Code merged into Intel official RDT (pqos) library
- **QEI: Query Acceleration Can be Generic and Efficient in the Cloud**
Y. Yuan, Y. Wang, R. Wang, R. Chowdhury, C. Tai, N. S. Kim
The IEEE International Symposium on High-Performance Computer Architecture (HPCA), 2021
- **Data Direct I/O Characterization for Future I/O System Exploration**
M. Alian, Y. Yuan, J. Zhang, R. Wang, M. Jung, N. S. Kim
The IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS), 2020
- **HALO: Accelerating Flow Classification for Scalable Packet Processing in NFV**
Y. Yuan, Y. Wang, R. Wang, J. Huang
The ACM/IEEE International Symposium on Computer Architecture (ISCA), 2019
- **Accelerating Distributed Reinforcement Learning with In-Switch Computing**
Y. Li, I. Liu, Y. Yuan, D. Chen, A. Schwing, J. Huang
The ACM/IEEE International Symposium on Computer Architecture (ISCA), 2019
- **Project Almanac: A Time-Traveling Solid-State Drive**
X. Wang, Y. Yuan, Y. Zhou, C. C. Coats, J. Huang
The ACM European Conference on Computer Systems (EuroSys), 2019
- **A Network-Centric Hardware/Algorithm Co-Design to Accelerate Distributed Training of Deep Neural Networks**
Y. Li, J. Park, M. Alian, Y. Yuan, Q. Zheng, P. Pan, R. Wang, A. Schwing, H. Esmaeilzadeh, N. S. Kim

Patents

- **Method and Apparatus for Scheduling Access to Multiple Accelerators**
R. Wang, **Y. Yuan**
US Patent App. 18/795,445, filed Aug. 2024
- **Method and Apparatus for Batching Pages for a Data Movement Accelerator**
R. Wang, **Y. Yuan**, R. Kuper
US Patent App. 18/477,628, filed Sep. 2023
- **Efficiently Merging Non-identical Pages in Kernel Same-page Merging (KSM) for Efficient and Improved Memory Deduplication and Security**
R. Kuper, R. Wang, **Y. Yuan**
US Patent App. 18/369,090, filed Sep. 2023
- **Data Consistency and Durability over Distributed Persistent Memory Systems**
R. Wang, **Y. Yuan**, Y. Wang, T.-Y. C. Tai, T. Hurson
US Patent 11,709,774, granted Jul. 2023
- **Workload Scheduler for Memory Allocation**
Y. Wang, R. Wang, T.-Y. C. Tai, **Y. Yuan**, P. Pathak, S. Vedantham, C. Macnamara
US Patent App. 16/799,745, filed Feb. 2020
- **Offload of Data Lookup Operations**
R. Wang, A. J. Herdrich, T.-Y. C. Tai, Y. Wang, R. Kondapalli, A. Bachmutsky, **Y. Yuan**
US Patent 11,698,929, granted Jul. 2023

Professional Services and Activities

- **Live Demo Presenter:** Improve System Performance by Offloading Memory-Intensive Kernel Features to CXL Type-2 Device (OCP Global Summit 2023)
- **Tutorial Organizer and Presenter:** On-chip Accelerators in 4th Gen Intel® Xeon® Scalable Processors: Features, Performance, Use Cases, and Future! (ISCA'2023)
- **Program Committee:** NSDI'2025, HPCA'2025, MICRO'2024 (ERC), ISCA'2024 (ERC and industry track committee), HPCA'2024, HPCA'2023 (ERC), EuroSys'2022 (shadow PC)
- **Reviewer:** IEEE Transactions on Parallel and Distributed Systems (TPDS, 2022), IEEE Computer Architecture Letter (CAL, 2022-2023), ACM Transactions on Architecture and Code Optimization (TACO, 2024)

Research Experience

- **Datacenter and System Taxes Reduction** 2023 – Present
Intel Labs and Meta
The concept of the “datacenter and systems taxes” refer to a set of shared low-level software components that consume a significant portion of processor cycles in warehouse-scale computers (WSCs) and distributed systems. We have been conducting various hardware and software innovations to reduce such taxes. The results have been published in *ISCA'24*, *ASPLOS'24*, *OCP 2023 Global Summit*, and *USENIX ATC'23*, and have been transferred to Intel (future) products.
- **Embracing Emerging CXL devices in Modern Datacenter** 2022 – Present
Intel Labs and Meta
CXL has been attracting much attention as the next generation of device interconnect standard, providing unique features such as memory expansion and cache coherence. **As the pioneer of this direction**, we have been exploring CXL memory devices (**Type 3**) and CXL accelerators (**Type 2**) based on real commodity hardware from different aspects, including both hardware functions enhancements and system software optimizations. The results have been published in *OSDI'24*, *VLDB'24*, *MICRO'23* and *OCP 2023 Global Summit*.
- **Accelerator Design for Network/Application Dataplane Operations** 2018 – 2023
UIUC and Intel Labs
Tackling the modern data explosion and the “killer microsecond” problem in datacenters, we design accelerator architecture, programming models, and integration schemes to accelerate a wide range of fine-grained but costly operations in datacenter’s software stacks and applications. The results have been published in *HPCA'23*, *HPCA'21* and *ISCA'19*.
- **I/O Subsystem Design and Optimization for Modern Server CPU** 2018 – 2021
UIUC and Intel Labs
High-speed I/O devices can exert significant pressure on the CPU’s cache/memory system. We study the I/O-host interaction behavior in the real system, and build realistic and accurate I/O subsystem models for gem5 simulator.

We also propose multiple solutions in both real systems and simulation models to optimize the data transfer, notification, and interference in the I/O subsystem. The results have been published in **CAL'23**, **MICRO'22**, **ISCA'21** and **ISPASS'20**.

• **In-network Computing for Distributed ML Training Acceleration**

2017 – 2021

UIUC and Microsoft Research

Distributed ML training is notoriously time- and resource-consuming. We propose to leverage the networking devices, including NICs (for in-network gradient compression) and switches (for in-network gradient aggregation), to facilitate the inter-machine communication, which is the most expensive portion in distributed training. We also explore the new potential for P4 programmable switch to process more complicated (floating-point) operations. The results have been published in **NSDI'22**, **ISCA'19**, and **MICRO'18**.

Teaching Experience

- **ECE 411:** Computer Organization and Design (UIUC, SP 2021)

Skills and Techniques

- **Programming languages:** C/C++, Verilog HDL, VHDL, Python, P4, Shell script, LaTeX, Matlab, etc.
- **Development skills:** Unix/Linux, FPGA, DPDK, RDMA, programmable switch, CUDA, gem5 simulator, sniper simulator, etc.